Data Science and the Data Scientist Toolkit

// FLATIRON SCHOOL

Agenda

- What is Data Science?
  - Roles and Responsibilities
  - The Process
- The Data Science Toolkit (Phase 1)

So:
What is
Data Science?

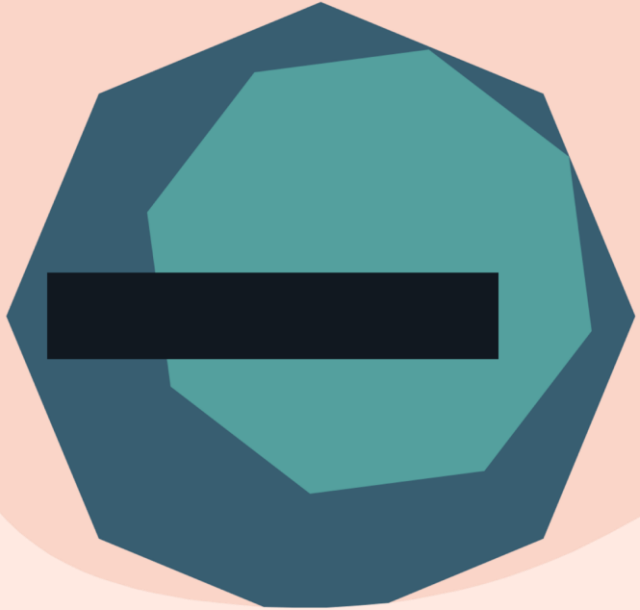

The Data Science Venn Diagram

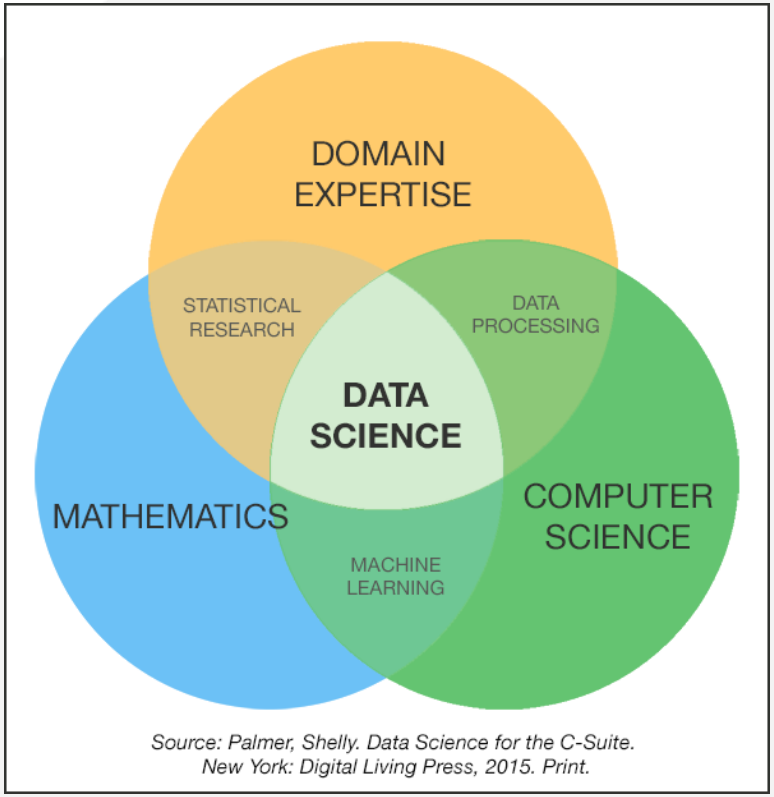

Source: Palmer, Shelly. Data Science for the C-Suite. New York: Digital Living Press, 2015. Print.

A data scientist is responsible for **collecting, analyzing and interpreting** data on various scales. **Offshoot of several traditional technical roles**, including mathematician, scientist, statistician and computer professional.

Common Roles & Responsibilities

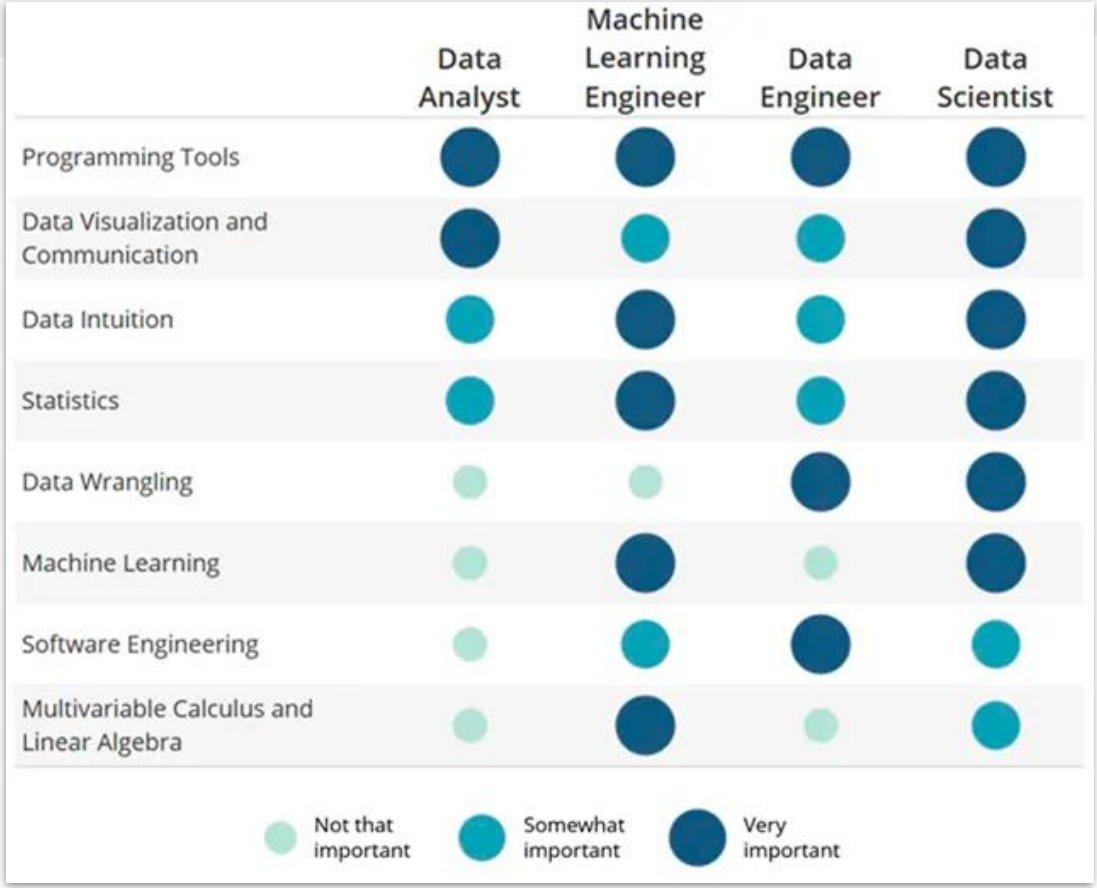

The Data Science Process

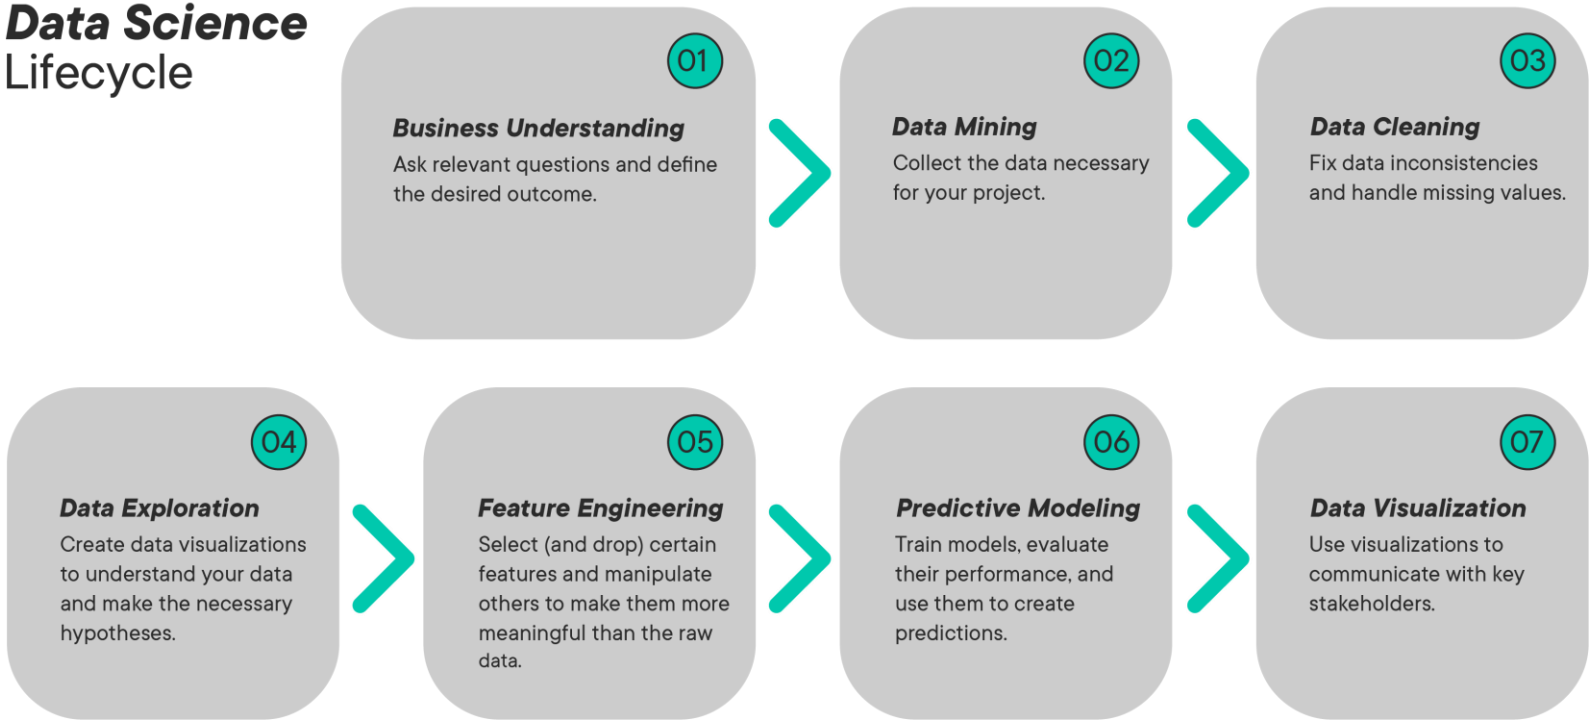

But it's actually an iterative process...

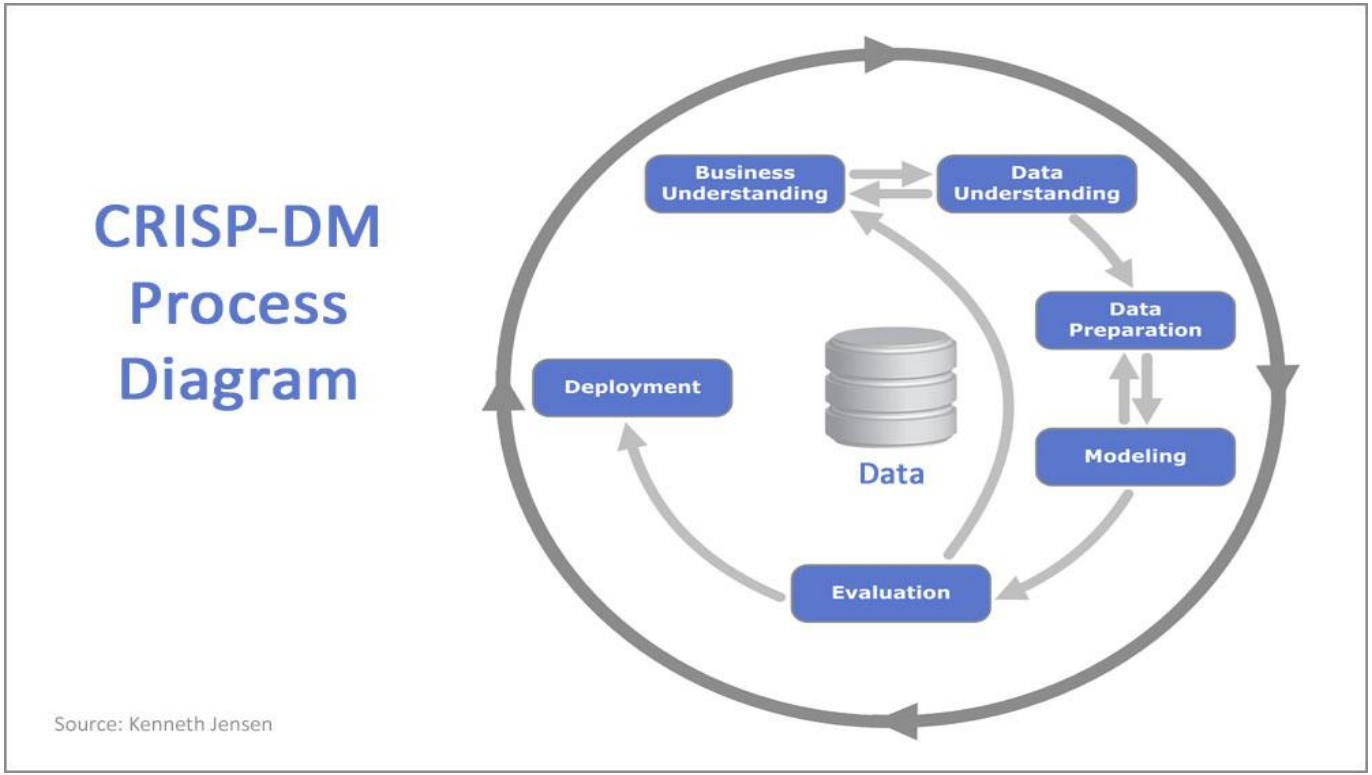

Asking the right questions

An irrelevant question + data/machine learning/stats = an irrelevant answer

Problem Formulation

Transformation:

Question into data science problem.

Some typical Data Science problems

Regression:

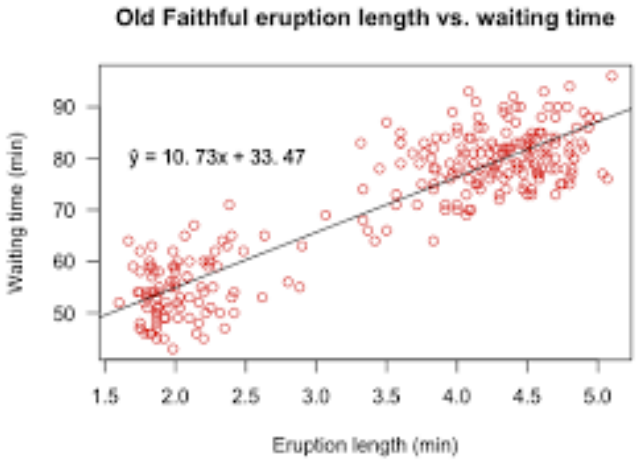

Old Faithful

Predict time between eruptions based on previous eruption duration.

Some typical Data Science problems

Classification:

Koala

or

Red Panda?

Some typical Data Science problems

Time series forecasting:

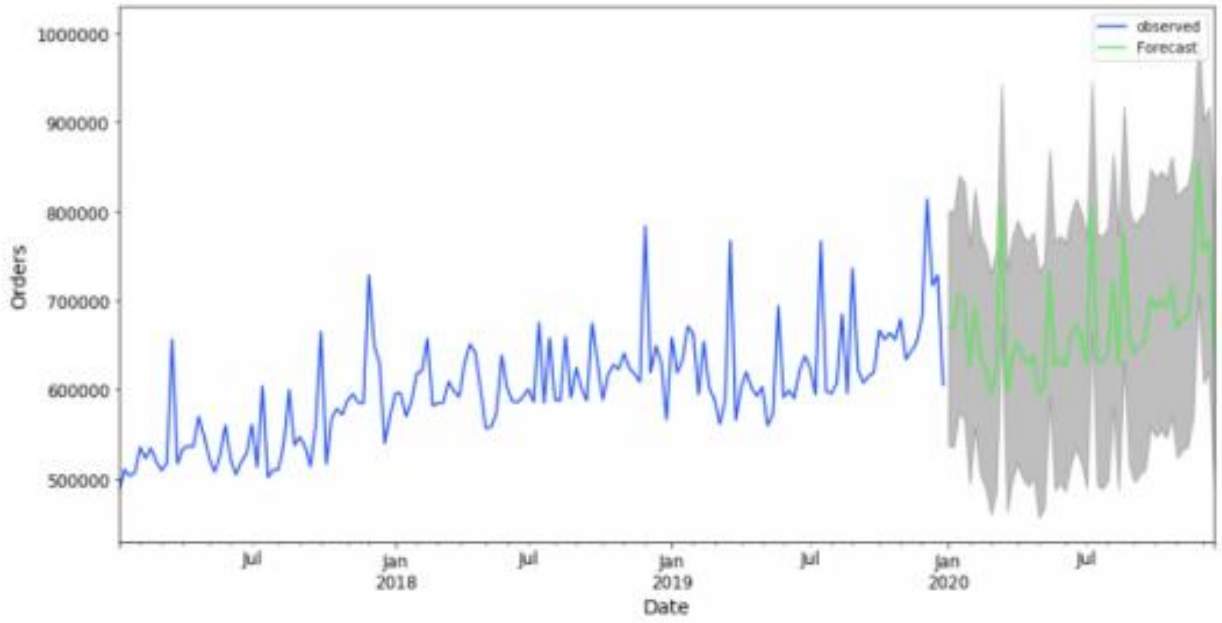

Some typical Data Science problems

Anomaly detection:

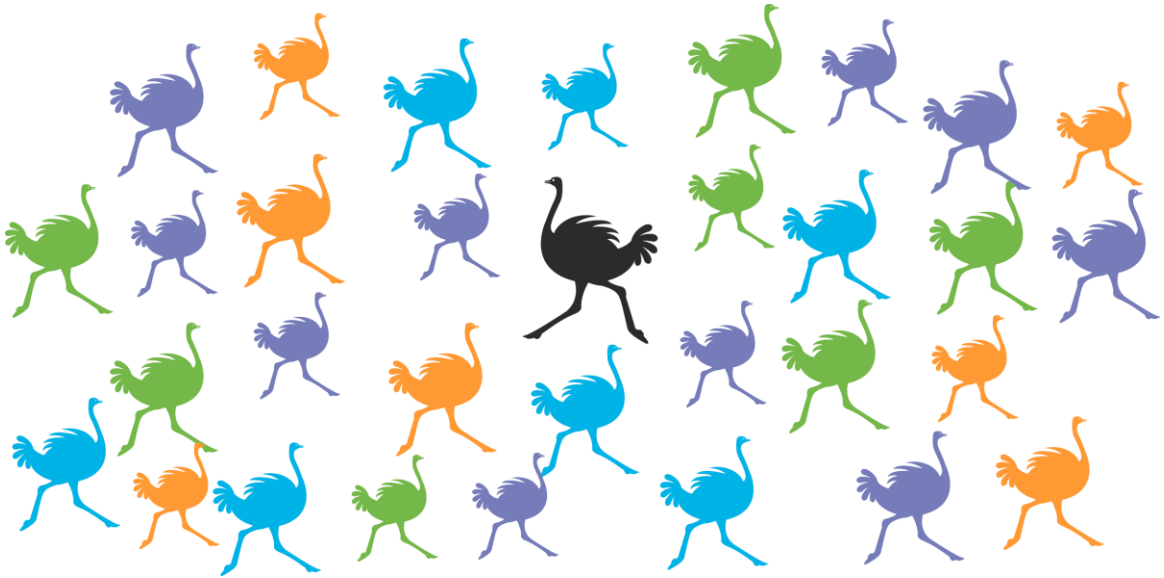

Data you might encounter

Be the data sculptor

Reshape the data:

- Clean and transform the data to your will.
- Data in useful form: modeling, answering your question.

An art form.

Where data scientists spend most of their time...

Data Cleaning Checklist

Up-to-date data
The data you use should be as recent as possible to ensure the maximum value of your results.

Missing values
Make sure to properly deal with missing values, as they may skew some of the results.

Duplicates
Check duplicates in your data and remove them as needed.

Outliers
Create a rule of thumb to spot outliers and remove them if needed.

Valid labels
Make sure to define valid labels for your categorical data.

Exploratory Data Analysis

Visualize and understand your data to transform into useful form.

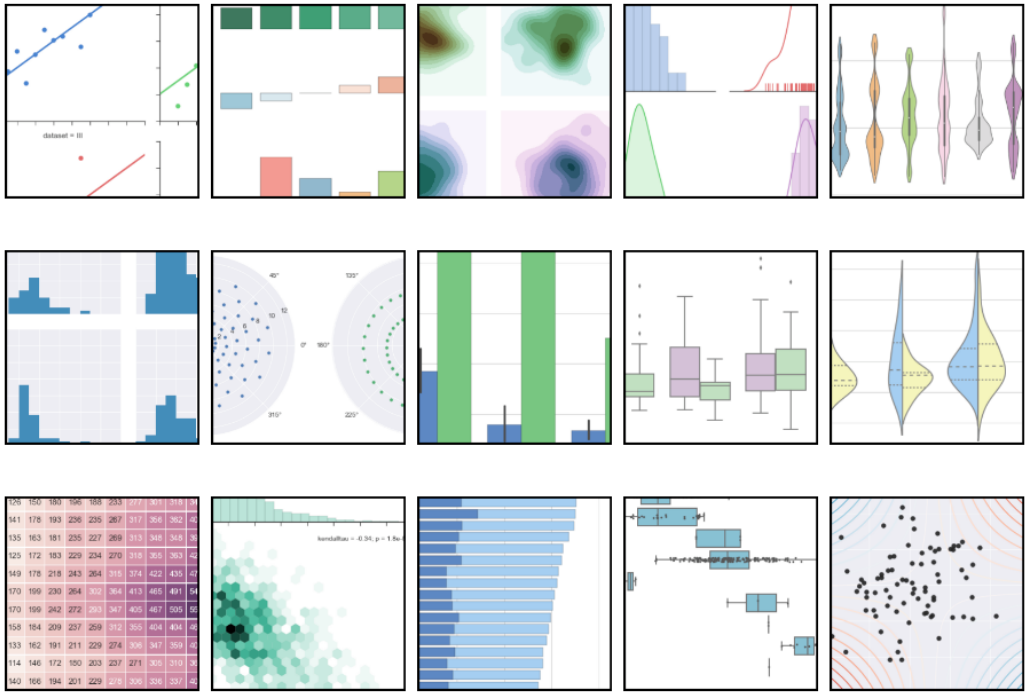

Feature Engineering

Transform raw data into meaningful features that directly address the problem you are trying to solve.

Modeling

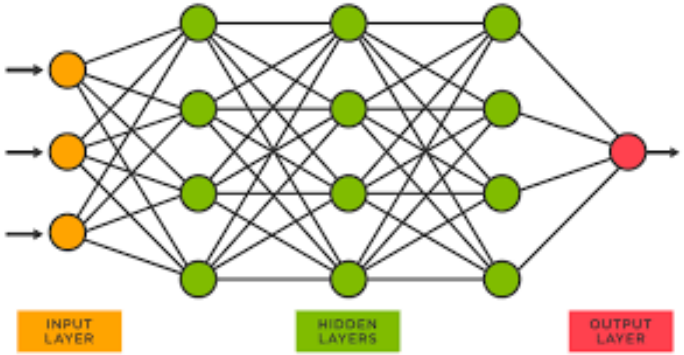

Neural Network.

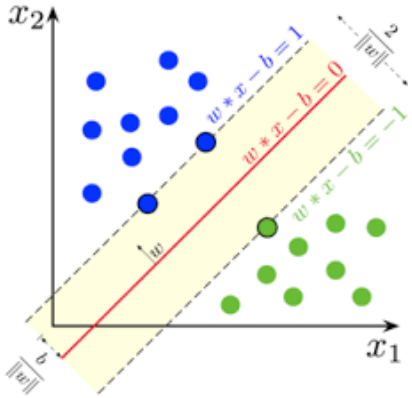

Support Vector Machine

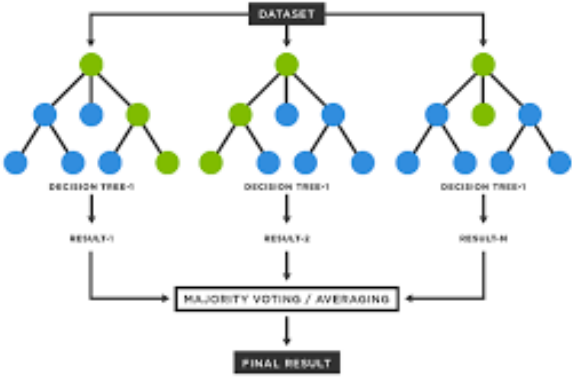

Random Forests

And more…try different models, tune, see what works best.

Presenting/visualizing results

This is key to a data scientist: Presentations/Reports

- **Know your audience**

- **State the problem clearly.**
- **How did you go about solving the problem?**
- **Key factors**
- **Visualizations of data and model**
- **Making recommendations.**

Bad Sign

Avoid

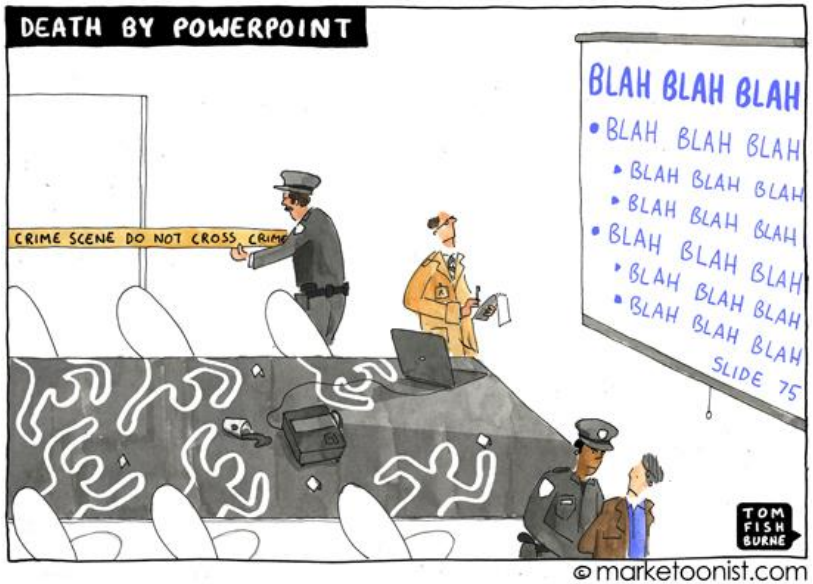

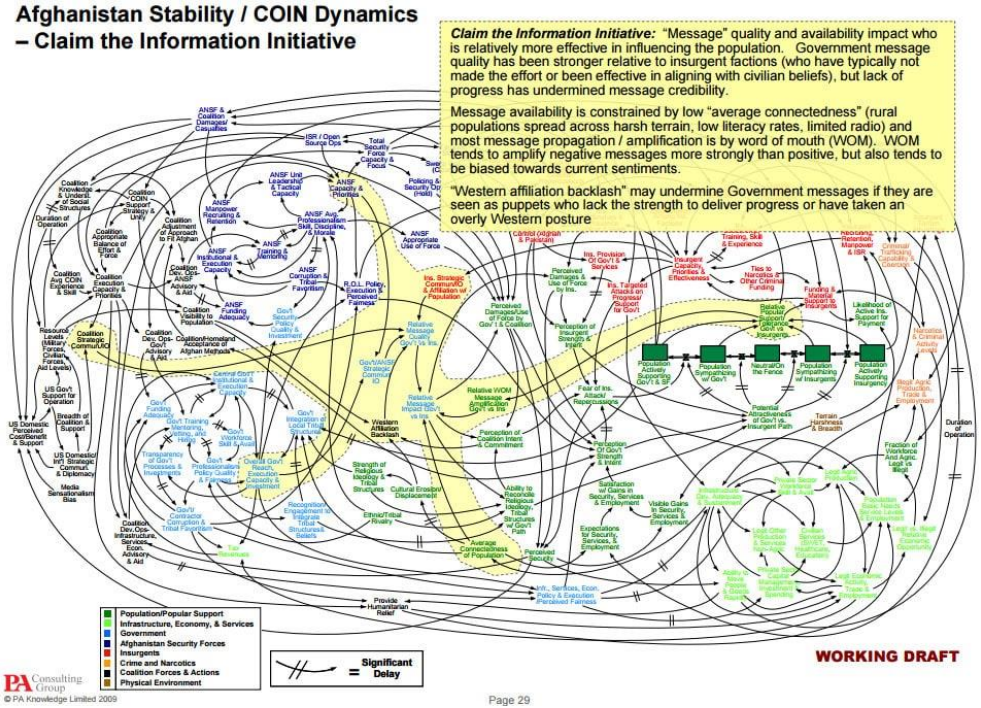

Yeah, OK.

Taste profile comparison: Islay Scotches

- **Ardbeg 10**: sweet, vanilla, lemon, lime, ardbeg, smoke, love, ridge, vanilla, mountain, peat, citrus, fruit, cloud, sea_spray, long, glorious, sea, caramel, beach_bonfire, smoke
- **Laphroaig 10**: 'seaweed', 'vanilla', 'ice_cream','tcp', 'plaster', 'oak', 'spice', 'cardamom', 'black_pepper', 'chilli', 'big', 'muscular', 'peat', 'spice', 'liquorice', 'big', 'dose', 'salt', 'slightly', 'sweet', 'beauty', 'classic', 'iodine', 'plaster', 'cool_wood', 'smoke', 'big', 'savoury', 'tarry', 'iodine'

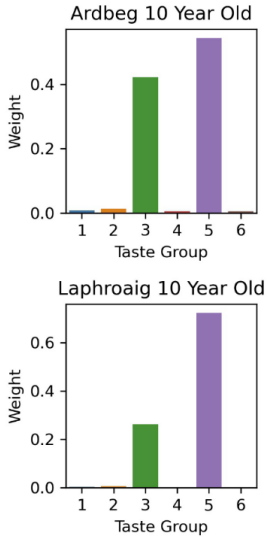

Figure: Taste group 3 = Herbal, tannin, citrus, wood spice. Drier notes. Taste group 5 = Peaty, salty, meaty notes.

Yeah, OK.

Taste profile comparison: Effect of Sherry Finish

- **Laphroaig 10**: 'seaweed', 'vanilla', 'ice_cream','tcp', 'plaster', 'oak', 'spice', 'cardamom', 'black_pepper', 'chilli', 'big', 'muscular', 'peat', 'spice', 'liquorice', 'big', 'dose', 'salt', 'slightly', 'sweet', 'beauty', 'classic', 'iodine', 'plaster', 'cool_wood', 'smoke', 'big', 'savoury', 'tarry', 'iodine'
- **Laphroaig 10 Sherry Finish**: 'roasted', 'cedar', 'peat_smoke', 'iodine', 'away', 'dark_chocolate', 'honey', 'vanilla_pod', 'meat', 'maple_syrup', 'bbq', 'lemon', 'charred_oak', 'smidge', 'coffee', 'balanced', 'finish', 'sherry', 'sweet', 'smouldering', 'peat'

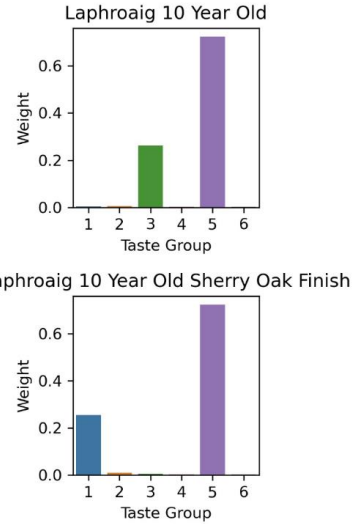

Figure: Taste group 3 = Herbal, tannin, citrus, wood spice. Drier notes. Taste group 5 = Peaty, salty, meaty notes. Taste group 1 = Nuts, molasses, candied berries, aromatic spice, and dark chocolate. Dark, sweet flavors

The Data Science Toolkit

Data Science Toolkit - Phase 1

Languages

Interfaces

Version Control

Versioning

Languages

Python

- Free, open source, versatile, powerful
- Not just for data science!
- Object-oriented (everything is an ‘object’)
- The Zen of Python

Structured Query Language (SQL)

- Connect to, change, and retrieve data from relational databases
- Developed in the 1970s, still going strong
- Many flavors

Interfaces

Jupyter Notebooks

- Streamlined document-centric interface for running and sharing code

IllumiDesk

- Hosts Jupyter Notebooks in the cloud

Code-Focused Text Editor

- Write text files in a code-native format
- **VS Code** is one of many that would work

Version Control

Git

- Distributed version tracking on any files
- Folder → “Repository”

GitHub

- Hosts Git repositories
- Collaborate and share code with others
- Backbone of the open source community
- Your Data Science portfolio!

Versioning

Anaconda

- Package management and deployment
- Designed with Data Science in mind
- Create and share environments

Python Package Index (PyPi)

- Database of public Python libraries
- Package installer (pip)
- Not everything is on Anaconda

Now:
Time to Get Started!